

# Introduction to programming for astronomers (fall 2016)

Thijs Coenen

Folkert Huizinga

September 7, 2017

# Contents

<b>I</b>	<b>Reading</b>	<b>3</b>
<b>1</b>	<b>Linux and its Command Line</b>	<b>4</b>
1.1	UNIX and Linux . . . . .	4
1.1.1	Users and Permissions . . . . .	4
1.1.2	The File System . . . . .	5
1.2	The Command Line . . . . .	5
1.2.1	Accessing the Shell . . . . .	5
1.2.2	Navigating the File System . . . . .	6
1.2.3	Running Programs . . . . .	8
1.2.4	Getting Help . . . . .	9
1.2.5	Copying, Moving and Removing Files . . . . .	9
1.2.6	Creating Directories . . . . .	12
1.2.7	Listing and Killing Processes . . . . .	12
1.2.8	Dealing with Text Files . . . . .	14
1.2.9	Glob Patterns . . . . .	16
1.2.10	Finding Files . . . . .	17
1.2.11	gzip and tar . . . . .	17
1.2.12	Permissions Continued . . . . .	18
1.2.13	Environment Variables . . . . .	20
1.2.14	UNIX Pipes . . . . .	20
1.2.15	Output Redirects . . . . .	21
<b>II</b>	<b>Assignments</b>	<b>23</b>
<b>2</b>	<b>Finding Pulsars from the command line</b>	<b>24</b>
2.1	Introduction . . . . .	24
2.1.1	Grading . . . . .	24
2.1.2	The assignments . . . . .	24

Part I

Reading

# Chapter 1

## Linux and its Command Line

In this chapter I will quickly introduce you to the very basics of using Linux through its command line interface. After a little history of Linux and operating systems like it, I will explain how you get access to the command line and go over some of the commands you will need. This chapter can be skipped safely if you are already familiar with the Linux command line. The examples in this chapter use the so called Bourne Again Shell (BASH). This shell comes standard with many flavors of Linux or UNIX.

### 1.1 UNIX and Linux

UNIX is an operating system developed in the early 1970s at AT&T Bell Labs. Unlike some of its contemporaries UNIX was designed as a multi user system from the start, it used a hierarchical file system (see 1.1.2) and consisted of many small programs that could be combined to perform complex operations. UNIX became popular in academia, where UNIX like systems are still used. In the early 1990s Linus Torvalds developed a new UNIX like operating system that he could run on his own PC. This operating system became known as Linux and, while initially developed by volunteers and hobbyists, it was soon picked up by businesses. Nowadays a large fraction of Linux' development is done by very large computer companies like Google, IBM and Samsung. Apple's Mac OS X is another popular UNIX version.

Currently Linux is used on computers ranging from mobile phones (e.g. Google's Android) to the largest super computer clusters. It furthermore runs a large fraction of the Internet's infrastructure (e.g. routers and web servers). Because UNIX was developed before graphical displays became generally available, its earliest interfaces were all text driven. Although graphical user interfaces are available for UNIX and Linux, textual interface persists. This textual user interface is also called a *command line* interface and on UNIX is provided through a so-called *shell*. Several different shells are available, but this chapter will only explain the basics of the Bourne Again Shell (BASH).

#### 1.1.1 Users and Permissions

Because UNIX was designed as a multi user operating system, it has a notion of users. Associated with every user are a level of access to the operating system. The so-called *root* user has full access to a UNIX system while the other users have more restricted access. On your own computer you will likely be able to log in as *root* while on shared computers you will only have access to your own files. Do note that even if you can log in as *root*, it is absolutely a bad idea to do so for day-to-day work. Since the *root* is all-powerful, a mistake while logged in as *root* can be much worse than a mistake while logged in as an ordinary user. Furthermore the software that you use may contain flaws or security problems that become a problem when that software (running as *root*) has full access to the computer. Furthermore each user is also assigned to a group of users that have the same system privileges.

UNIX controls access to files and directories based on ownership and so-called *permissions*. Each file and directory on a UNIX system has an *owner* (one of the users on that system) and permissions. UNIX has three permissions: *read*, *write* and *execute*. Read permission is needed to be able to view the contents of that file. Write permission is needed on directories to create files in that directory. For files write permission is needed to change or erase them. Execute permission is needed on directories to be able to view their contents and on files to execute them. Permission may also be set at the *group* level or for all *other* users with access to a computer system (see Section 1.2.12).

### 1.1.2 The File System

UNIX uses a hierarchical file system meaning that directories in it can be arbitrarily nested. Each directory can contain files or sub directories — the former directory is also referred to as *parent* while the latter are referred to as *children*. The directory hierarchy is a tree with the directory at the base of that tree referred to as *root*, denoted as `/`.

When specifying locations, or *paths*, in a file system two different notations may be used. *Absolute paths*, on the one hand, describe the location of a file or directory in absolute terms, i.e. without reference to some other location on the file system. *Relative paths*, on the other hand, specify a location in the file system relative to some other location in the file system (usually the current location of the user). Absolute paths start at the root, i.e. they start with `/`, while relative paths do not.

Because a UNIX system may be shared between many users, each user has his or her own directory to store files in. This special directory is called a *home directory*. Home directories can be found in `/home` on Linux systems and in `/Users` for Mac OS X systems. E.g. my user (`thijscoenen`) has the home directory `/Users/thijscoenen` on a Mac OS X system and `/home/thijscoenen` on a Linux system. The UNIX(-like) operating systems have very specific file system layouts, with some directories assigned to programs or even parts of the operating system itself. Most directories that reside directly in the root directory are in fact system directories and as a normal user you will rarely need access to them. Below, as an example, the contents of my Mac OS X machine's root directory none of which is part of my own files or directories<sup>1</sup>

```
Gretchen:coolstuff thijscoenen$ ls /
Applications  System      cores      mach_kernel  tmp
Developer     Users       dev        net          usr
Library        Volumes    etc        private      var
Network        bin        home       sbin
```

Gretchen:coolstuff thijscoenen\$

UNIX also allows so-called *hidden files* that are normally invisible. Any filename or directory name that starts with a dot will be hidden. Hidden files are generally used for settings or for those files that a user does not need access to directly (only through some program). E.g. the settings for the BASH shell are kept in a file called `.bash.profile` or `.bashrc` in your home directory.

## 1.2 The Command Line

### 1.2.1 Accessing the Shell

The command line interface, provided by a shell program, is generally accessible through a terminal emulator program. The name terminal derives from the simple computers (terminals) that were used in the past to access large shared computer systems. When you start a terminal emulator you are generally dropped into your home directory. On Mac OS X you can use the “Terminal” program, while on a Linux systems with KDE you use “Konsole”, on Linux systems with Gnome “GNOME Terminal” and on Linux systems with Unity the program is also called “Terminal”. This

<sup>1</sup>The example shows some directories that are Mac OS X specific: Applications, Developer, Library, Network, System and Volumes.




Figure 1.1: A screenshot showing Mac OS X’s built-in Terminal program. This program allows you to access the underlying BASH shell on Mac OS X.

chapter will only explain the basics of the BASH shell as it is standard on many Linux systems and on all recent versions of Mac OS X.

## 1.2.2 Navigating the File System

Assuming you have access to the command line, you will be presented with a so called prompt. What the prompt looks like depends on the version of UNIX or Linux you are using and how it was configured. In my case the command prompt looks like<sup>2</sup>:

```
Gretchen:~ thijscoenen$
```

You can issue your commands by typing them followed by a press of the return key. One of the first things you may ask yourself, is where am I? Usually when you start a new command line session you will be dropped off in your user’s home directory. To find out where in the file system you are you use the `pwd` (print working directory) command. The following example shows you what the output looks like for me:

```
Gretchen:~ thijscoenen$ pwd
/Users/thijscoenen
Gretchen:~ thijscoenen$
```

The home directory in these examples has a subdirectory called `coolstuff` that we will access. To change directories to the `coolstuff` directory use the `cd` (change directory) command. In the following example a `pwd` command issued as well to show we have actually changed directory:

---

<sup>2</sup>What your prompt looks like exactly depends on the settings of your shell and is a matter of taste. My computer shows `computername:directory username$`, with Gretchen being the name of my computer.

```
Gretchen:~ thijscoenen$ cd coolstuff
Gretchen:coolstuff thijscoenen$ pwd
/Users/thijscoenen/coolstuff
Gretchen:coolstuff thijscoenen$
```

The eagle eyed among you may have noticed the `~` (tilde) appearing in the shell session examples. This is shorthand for your home directory, and your shell will expand it to your full home directory. You can use the tilde in commands to shorten them. To return to your home directory use the `cd ~` command.

What does the `coolstuff` directory actually contain? To get a listing of the contents of a directory you can use the `ls` command. This command takes many options but the easiest way of using it is as follows:

```
Gretchen:coolstuff thijscoenen$ ls
alice.txt      hello          noperm
Gretchen:coolstuff thijscoenen$
```

Hidden files on UNIX systems have names that start with a dot. There are two special directories that are always present in every directory, the single dot and double dot directories. The former `.` is shorthand for the current directory and the latter `..` is shorthand for the parent directory. By passing the `-a` option to `ls` you make hidden files visible.

```
Gretchen:coolstuff thijscoenen$ ls -a
.      ..      alice.txt      hello          noperm
Gretchen:coolstuff thijscoenen$ cd ..
Gretchen:~ thijscoenen$ pwd
/Users/thijscoenen
Gretchen:~ thijscoenen$
```

The `ls` command has many options, but several are worthy of mention. First, the `-l` option will list the directory contents with one file or directory per line and also show the size (in bytes) and permissions for each entry. Second, the `-lh` option will show the file sizes in human readable format (so not in bytes for large files, see below for an example).

```
Gretchen:coolstuff thijscoenen$ ls -l
total 312
-rw-r--r--  1 thijscoenen  staff  147731 22 jul 23:27 alice.txt
-rwxr--r--  1 thijscoenen  staff     43  7 aug 18:02 hello
-rw-r--r--  1 thijscoenen  staff     50 11 aug 14:35 noperm
Gretchen:coolstuff thijscoenen$ ls -lh
total 312
-rw-r--r--  1 thijscoenen  staff   144K 22 jul 23:27 alice.txt
-rwxr--r--  1 thijscoenen  staff    43B  7 aug 18:02 hello
-rw-r--r--  1 thijscoenen  staff    50B 11 aug 14:35 noperm
Gretchen:coolstuff thijscoenen$
```

Another nice command is `ls -lrt` which will list all files and sorted by the time of last access (oldest to newest). This is handy when you want to find the most recently accessed file in a directory as it will be at the end of the `ls` output.

The `ls` command can furthermore be used to list the content of a specific directory. Using the previous example I can list the contents of the `coolstuff` directory from my home directory using the command `ls coolstuff`. When you want to recursively list the contents of some directory you can use the `ls -R` command. It shows the content of the current directory, the content of the subdirectories and the content of the subdirectories of each sub directory etc. Because this command produces a lot of output I suggest you try it for yourself.

### 1.2.3 Running Programs

The shell can only start programs it can find, and the directories where the shell will look for programs are listed in an environment variable (see Section 1.2.13 for an explanation) called `$PATH`. For programs that are on the `$PATH`, i.e. in one of the directories listed by the `$PATH` variable, you can just type their name and press enter (even if the program is in a different directory than you are). The programs `ls` and `cd` are on the `$PATH` (on my system they actually reside in the `/bin` directory). When you try to run a non-existent program or a program that your shell cannot locate, you will see an error like the following:

```
Gretchen:coolstuff thijscoenen$ ls
alice.txt      hello          noperm
Gretchen:coolstuff thijscoenen$ nosuchthing
-bash: nosuchthing: command not found
Gretchen:coolstuff thijscoenen$
```

This will even happen for programs present in the same directory as you — unless that directory happens to be on the `$PATH`. To remedy this problem, prepend the program name with `./` (this tells your shell to look in the current directory).

```
Gretchen:coolstuff thijscoenen$ hello
-bash: hello: command not found
Gretchen:coolstuff thijscoenen$ ./hello
Hello world!
Gretchen:coolstuff thijscoenen$
```

You may not have sufficient permissions to run just any program on UNIX. To check whether you have permissions to run the program you can just attempt to run it, your shell will report an error if you have insufficient permissions.

```
Gretchen:coolstuff thijscoenen$ ./noperm
-bash: ./noperm: Permission denied
Gretchen:coolstuff thijscoenen$
```

All the previous examples were for programs that run for a short time, if you have long running programs you may want to continue issuing new commands without waiting for some program to finish. You could just start a second (or third etc.) terminal and continue issuing commands in that new terminal. The BASH shell however allows you to start programs in the background so that it is possible to keep issuing commands in that same shell. To start a program in the background just add a space and ampersand `&` to the command. You can move a program from the background to the foreground with the `fg` command:

```
Gretchen:background-demonstration thijscoenen$ ls
runs-1-minute
Gretchen:background-demonstration thijscoenen$ ./runs-1-minute &
[1] 55868
Gretchen:background-demonstration thijscoenen$ pwd
/Users/thijscoenen/background-demonstration
Gretchen:background-demonstration thijscoenen$ fg
./runs-1-minute
One minute passed
Gretchen:background-demonstration thijscoenen$
```

As you can see in the example above a number is shown in the shell after the program was started using the `&`. This number is the process number of the program just started, in Section 1.2.7 these process numbers will be explained. If you started a program in the foreground, but want to move it to the background you first suspend the program and then move it to the background. Suspend the program by pressing `[Control] + [Z]` and then move it to the background by issuing the `bg` command.



```
Gretchen:background-demonstration thijscoenen$ ./runs-1-minute
^Z
[1]+  Stopped                  ./runs-1-minute
Gretchen:background-demonstration thijscoenen$ bg
[1]+ ./runs-1-minute &
Gretchen:background-demonstration thijscoenen$ pwd
/Users/thijscoenen/background-demonstration
Gretchen:background-demonstration thijscoenen$ fg
./runs-1-minute
One minute passed
Gretchen:background-demonstration thijscoenen$
```

In this example you can see that the `[Control] + [Z]` key presses are shown in the shell as `^Z`. It is also possible to terminate a program running in the shell by pressing `[Control] + [C]`. Some programs may not react to an attempt to terminate it this way and you may have to resort to the `kill` command described below in Section 1.2.7.

### 1.2.4 Getting Help


While this reader will get you started using Linux it cannot possibly explain all features of BASH or all programs available by default on a UNIX machine. Fortunately a large amount of documentation is available, some of it accessible directly from the command line and some of it on line. Most command line programs have simple built in help that is usually displayed when no options are specified on the command line or through options as `-?`, `-h` or `--help`.

While many programs provide some simple documentation about themselves UNIX provides a standard command line documentation viewer called `man`. To get help about a specific program just type the command `man program-name`, where you should replace `program-name` with the program that you want help for. Figure 1.2 shows the help for `ls` (accessed by typing `man ls`). The documentation can be navigated using the `[up]` and `[down]` arrows, to move up and down by a line, and `[space]` to jump a full page. The documentation can also be searched by typing a forward slash `/` followed by the word you are looking for and then `[return]`. When there are several matches for a search you can jump to the next match by pressing `[n]` (or in reverse with `[N]`). The viewer can be quit by typing a `[q]`. More information about navigating in `man` can be accessed by pressing `[h]` while it is running. The documentation available through `man` are also called “man pages”. Some software will be documented using `info`, try running that like you would `man` in case no man pages are available (because the documentation may be written in `info` format).

On line there are several good resources for help with UNIX problems, that can be very helpful because the man pages are quite technical and at times hard to read. The very basic UNIX commands tend to have Wikipedia pages with examples. Stack overflow, a community site about programming problems (see <http://stackoverflow.com>), can be helpful as can be the related Stack Exchange site about UNIX and Linux (see <http://unix.stackexchange.com>). A nice site that can explain some shell commands is Explain Shell (see <http://explainshell.com>), it allows you to cut and paste a command and it then shows you an explanation of it. The Linux Documentation Project (see <http://www.tldp.org>) is also useful. A note about using Google to find documentation on the Internet: not everything you find will be correct so refrain from just copying whatever you find!

### 1.2.5 Copying, Moving and Removing Files

To copy files, or directories use the `cp` command. The general shape of the command is `cp source destination`. If the source is a file and the destination is also a filename the contents of `source` will be copied over the destination file. If `destination` does not exist yet `source` will be copied to a file with the specified destination as name. If `destination` is a directory the `source` file will be copied to it. If the source is a directory itself the copying will fail (these examples use `cat` to show file contents):

Figure 1.2: The man documentation viewer showing part of the documentation for `ls`.

```
Gretchen:cp-demonstration thijscoenen$ ls -l
total 16
drwxr-xr-x  2 thijscoenen  staff   68 21 aug 22:40 directory_a
-rw-r--r--  1 thijscoenen  staff   16 21 aug 22:40 file_a.txt
-rw-r--r--  1 thijscoenen  staff   16 21 aug 22:40 file_b.txt
Gretchen:cp-demonstration thijscoenen$ cat file_a.txt
This is file A.
Gretchen:cp-demonstration thijscoenen$ cat file_b.txt
This is file B.
Gretchen:cp-demonstration thijscoenen$ cp file_a.txt file_c.txt
Gretchen:cp-demonstration thijscoenen$ ls -l
total 24
drwxr-xr-x  2 thijscoenen  staff   68 21 aug 22:40 directory_a
-rw-r--r--  1 thijscoenen  staff   16 21 aug 22:40 file_a.txt
-rw-r--r--  1 thijscoenen  staff   16 21 aug 22:40 file_b.txt
-rw-r--r--  1 thijscoenen  staff   16 21 aug 22:41 file_c.txt
Gretchen:cp-demonstration thijscoenen$ cat file_c.txt
This is file A.
Gretchen:cp-demonstration thijscoenen$ cp file_a.txt file_b.txt
Gretchen:cp-demonstration thijscoenen$ cat file_b.txt
This is file A.
Gretchen:cp-demonstration thijscoenen$ cp file_a.txt directory_a/
Gretchen:cp-demonstration thijscoenen$ ls directory_a/
file_a.txt
Gretchen:cp-demonstration thijscoenen$ cp directory_a whatever
cp: directory_a is a directory (not copied).
```

```
Gretchen:cp-demonstration thijscoenen$
```

To copy a directory and all it contains use the `-r` option of `cp`:

```
Gretchen:cp-demonstration thijscoenen$ cp -r directory_a/ whatever
Gretchen:cp-demonstration thijscoenen$ ls whatever/
file_a.txt
Gretchen:cp-demonstration thijscoenen$
```

To just move a file to a different filename or location use the `mv` command. Simple use of `mv` takes the shape `mv source destination`. If `source` is a file and `destination` does not exist yet `source` will be moved there (after this operation `source` will cease to exist). If `destination` does exist, and is a file, it will be overwritten. If `destination` is a directory the file `source` will be moved to that directory. If `source` itself is a directory it will be moved to `destination`, with similar rules:

```
Gretchen:cp-demonstration thijscoenen$ ls -l
total 24
drwxr-xr-x  3 thijscoenen  staff  102 21 aug 22:42 directory_a
-rw-r--r--  1 thijscoenen  staff   16 21 aug 22:40 file_a.txt
-rw-r--r--  1 thijscoenen  staff   16 21 aug 22:42 file_b.txt
-rw-r--r--  1 thijscoenen  staff   16 21 aug 22:41 file_c.txt
drwxr-xr-x  3 thijscoenen  staff  102 21 aug 22:45 whatever
Gretchen:cp-demonstration thijscoenen$ mv file_a.txt file_b.txt
Gretchen:cp-demonstration thijscoenen$ ls -l
total 16
drwxr-xr-x  3 thijscoenen  staff  102 21 aug 22:42 directory_a
-rw-r--r--  1 thijscoenen  staff   16 21 aug 22:40 file_b.txt
-rw-r--r--  1 thijscoenen  staff   16 21 aug 22:41 file_c.txt
drwxr-xr-x  3 thijscoenen  staff  102 21 aug 22:45 whatever
Gretchen:cp-demonstration thijscoenen$ mv whatever directory_a/
Gretchen:cp-demonstration thijscoenen$ ls -l
total 16
drwxr-xr-x  4 thijscoenen  staff  136 21 aug 22:57 directory_a
-rw-r--r--  1 thijscoenen  staff   16 21 aug 22:40 file_b.txt
-rw-r--r--  1 thijscoenen  staff   16 21 aug 22:41 file_c.txt
Gretchen:cp-demonstration thijscoenen$ mv directory_a/ file_b.txt
mv: rename directory_a/ to file_b.txt: Not a directory
Gretchen:cp-demonstration thijscoenen$
```

As you can see trying to move a directory to an existing file will fail.

Removing files is done with the `rm` command. Generally you just use the `rm somefile` command. If `somefile` is a directory this will fail because only *empty* directories are removed with `rmdir`.

```
Gretchen:cp-demonstration thijscoenen$ ls -l
total 16
drwxr-xr-x  4 thijscoenen  staff  136 21 aug 22:57 directory_a
-rw-r--r--  1 thijscoenen  staff   16 21 aug 22:40 file_b.txt
-rw-r--r--  1 thijscoenen  staff   16 21 aug 22:41 file_c.txt
Gretchen:cp-demonstration thijscoenen$ rm file_b.txt
Gretchen:cp-demonstration thijscoenen$ ls -l
total 8
drwxr-xr-x  4 thijscoenen  staff  136 21 aug 22:57 directory_a
-rw-r--r--  1 thijscoenen  staff   16 21 aug 22:41 file_c.txt
Gretchen:cp-demonstration thijscoenen$ rm directory_a/
rm: directory_a/: is a directory
```

```
Gretchen:cp-demonstration thijscoenen$ rmdir directory_a/
rmdir: directory_a/: Directory not empty
Gretchen:cp-demonstration thijscoenen$
```

If you want to remove a directory and all its content (whether files or sub directories) you can use the `rm -rf *`. Note though that this is a very dangerous command, if not the “most dangerous command ever”<sup>3</sup>, as it will recursively (the `-r` option) remove everything (matched by the `*` “glob pattern”) and skip any questions (the `-f` option). For more information on glob patterns see Section 1.2.9. Note that `rm -rf *` cannot be undone!

## 1.2.6 Creating Directories

Because in the previous section it was shown how to remove directories you may ask yourself how you create them in the first place. UNIX has the `mkdir` command for that. You can create a directory with the command `mkdir somedirectory` where `somedirectory` is the name of the directory to be created. It is also possible to create several directories in one go by listing their names after `mkdir` separated with spaces. If you have to create a directory inside of several that is possible using the command `mkdir -p` followed the full path you want to create.

```
Gretchen:mkdir-examples thijscoenen$ ls
Gretchen:mkdir-examples thijscoenen$ mkdir dir1 dir2 dir3
Gretchen:mkdir-examples thijscoenen$ ls
dir1      dir2      dir3
Gretchen:mkdir-examples thijscoenen$ ls -l
total 0
drwxr-xr-x  2 thijscoenen  staff  68 21 aug 23:28 dir1
drwxr-xr-x  2 thijscoenen  staff  68 21 aug 23:28 dir2
drwxr-xr-x  2 thijscoenen  staff  68 21 aug 23:28 dir3
Gretchen:mkdir-examples thijscoenen$ mkdir -p container/subdir/subdir/whatever
Gretchen:mkdir-examples thijscoenen$ cd container/subdir/subdir/whatever
Gretchen:whatever thijscoenen$ pwd
/Users/thijscoenen/mkdir-examples/container/subdir/subdir/whatever
Gretchen:whatever thijscoenen$
```

## 1.2.7 Listing and Killing Processes

On UNIX each running program consists of one or sometimes more processes that are identified by *process numbers* often abbreviated to PID. To get an idea of the running programs use the `top` command. This will show you a continually updating list of running processes, sortable by for instance memory use or processor use. When your computer is stuck it is quite often possible to find the offending program by looking through the most active processes in `top`. As with `man` you can exit `top` by pressing `q`. Figure 1.3 shows a screenshot of `top`. You can get a list of the currently running processes using the `ps` command, which unlike `top` will not continually update - it is like an `ls` for processes. Run without options `ps` only shows the processes running in the current shell while you can get a list of all processes running in all shells for a certain user using `ps -u username` (with `username` replaced with relevant username). For example, while logged in on a LOFAR<sup>4</sup> compute node as `coenen`:

```
coenen@locus048:~$ ps
  PID TTY          TIME CMD
 14099 pts/2    00:00:00 bash
 15518 pts/2    00:00:00 ps
```

<sup>3</sup>According to one of my proofreading colleagues.

<sup>4</sup>The Low Frequency Array, a large radio telescope array operating at low radio frequencies that has its core in the Netherlands.

```

Terminal — top — 80x24
Processes: 62 total, 2 running, 60 sleeping, 267 threads          17:29:29
Load Avg: 0.13, 0.20, 0.17  CPU usage: 4.24% user, 9.43% sys, 86.32% idle
SharedLibs: 5196K resident, 5648K data, 0B linkedit.
MemRegions: 13931 total, 632M resident, 13M private, 228M shared.
PhysMem: 685M wired, 947M active, 311M inactive, 1942M used, 105M free.
VM: 156G vsize, 1039M framework vsize, 17259915(0) pageins, 6093270(0) pageouts.
Networks: packets: 64936199/83G in, 37656757/3714M out.
Disks: 28389042/396G read, 17131423/460G written.

  PID  COMMAND      %CPU  TIME    #TH  #WQ  #PORT  #MREG  RPRVT  RSHRD  RSIZE
58153  screencaptur  0.4   00:00.03  3    2    42    81    232K-  9000K+ 2812K
58151  top           5.8   00:00.33  1/1   0    25    33    984K+  264K   1564K+
58148  bash          0.0   00:00.00  1     0    17    24    392K   744K   1064K
58147  login         0.0   00:00.01  1     0    22    53    496K   312K   1596K
57871  vim           0.0   00:25.76  1     0    17    36    1692K  244K   3028K
57863  bash          0.0   00:00.01  1     0    17    24    408K   744K   1108K
57862  login         0.0   00:00.01  1     0    22    53    488K   312K   1596K
57789  bash          0.0   00:00.03  1     0    17    24    416K   744K   1116K
57788  login         0.0   00:00.01  1     0    22    53    488K   312K   1596K
57747  bash          0.0   00:00.04  1     0    17    24    424K   744K   1116K
57746  login         0.0   00:00.01  1     0    22    53    488K   312K   1596K
57729  bash          0.0   00:00.01  1     0    17    24    396K   744K   1064K
57728  login         0.0   00:00.01  1     0    22    53    488K   312K   1596K
57706  Preview       0.0   00:13.14  2     1   119   277    16M+   27M-   35M

```

Figure 1.3: This screenshot shows top running on a Mac OS X system.

```

coenen@locus048:~$ ps -u coenen
  PID TTY          TIME CMD
14098 ?           00:00:00 sshd
14099 pts/2       00:00:00 bash
15229 ?           00:00:00 sshd
15230 pts/3       00:00:00 bash
15266 pts/3       00:00:00 top
15522 pts/2       00:00:00 ps
coenen@locus048:~$

```

Another useful option to `ps` is `-A` which will select all processes, also not necessarily running in a terminal.

The process numbers can be used to terminate unresponsive programs or those that are otherwise stuck. The `kill` command allows you to terminate processes by process number:

```

Gretchen:background-demonstration thijscoenen$ ./runs-1-minute &
[1] 59380
Gretchen:background-demonstration thijscoenen$ kill 59380
Gretchen:background-demonstration thijscoenen$ fg
-bash: fg: job has terminated
[1]+  Terminated                  ./runs-1-minute
Gretchen:background-demonstration thijscoenen$

```

Some processes may still not terminate, you can then send them a stronger message using the `-9` (KILL) signal. E.g. to kill process 101 using the KILL signal: `kill -9 101`.




Figure 1.4: The `less` program showing a text file, you can move through the file using the arrow keys.

### 1.2.8 Dealing with Text Files

You will inevitably have to work with text files, because, among other reasons, programming source code are text files, many of UNIX's settings are in small text files in your home directory and quite a few scientific data sets are encoded in text files. Luckily UNIX has many tools to work with text files from the command line. The first command is `cat` which will show the contents of a text file in your terminal. Since `cat` does not paginate, this will likely cause your terminal to scroll. In the following example I look at the contents of the `hello` file (which is a Python script and therefore text):

```
Gretchen:coolstuff thijscoenen$ ls
Whatever      alice.txt      hello          noperm
Gretchen:coolstuff thijscoenen$ cat hello
#!/usr/bin/env python
print "Hello world!"
Gretchen:coolstuff thijscoenen$
```

If the file you are working with is large, then using `cat` to look at its content is inconvenient. By using the `more` or `less` commands you will be presented with a so-called pager that allows you to view and move around in the text file without scrolling. The `less` command allows you to move in both directions through the file using the up and down arrow keys. You can search for words by pressing a forward slash `/` followed by the word you are looking for and then `enter`. If there are several matches you can move between them using `n` (or `N` for the reverse direction). Help is accessible by pressing `h` and you can quit `less` by pressing `q`. This is very similar to the `man` works because it actually uses `less` to show the actual man page (visible in the title bar of the window in Figure 1.2).

If you are only interested in a quick peek at the contents of some text file, you can use the `head` or `tail` commands to, respectively, look at the first few lines or last few lines of a text file. Similarly if you want to find out the number of words or lines in a file the `wc` (word count) command is useful, without options it shows the number of lines, the number of words and the number of bytes in the file.

```
Gretchen:coolstuff thijscoenen$ head alice.txt
ALICE'S ADVENTURES IN WONDERLAND
```

```
Lewis Carroll
```

```
CHAPTER I. Down the Rabbit-Hole
```

```
Alice was beginning to get very tired of sitting by her sister on the
bank, and of having nothing to do: once or twice she had peeped into the
book her sister was reading, but it had no pictures or conversations in
```

```
Gretchen:coolstuff thijscoenen$
```

```
Gretchen:coolstuff thijscoenen$ wc alice.txt
```

```
3335  26444 147731 alice.txt
```

```
Gretchen:coolstuff thijscoenen$ wc -l alice.txt
```

```
3335 alice.txt
```

```
Gretchen:coolstuff thijscoenen$
```

The last example (the `wc -l alice.txt` command) shows you how to count only the number of lines in a file. A oft used file format to encode tabular data is a so called Comma Separated Values (CSV) file (white space or tab separation are also used frequently). The lines in these files generally correspond to one data point, you can thus get a quick feel for the size of the data set by using `wc -l` on that file.

So far we have only treated looking through files, not editing or creating them. Many different text editors are available on UNIX, spanning the gamut from editors that run in your terminal and use arcane commands to more modern text editors with a graphical user interface. Note though that in this context modern does not always mean more efficient! The choice of editor is personal and there is not enough space in this reader to explain how to use them all. Instead, I have added a chapter to the appendix that lists some good text editors (see Appendix ??). For now I will just mention that you will want a text editor that is focused on programmers as they allow you to work more efficiently.

To search a file for a certain piece of text you can use the `grep` command. To find out on which lines the Cheshire cat appears in the text we use `grep` with the `-n` option to print line numbers and `-i` option to perform case insensitive matches:

```
Gretchen:coolstuff thijscoenen$ grep -ni "Cheshire" alice.txt
```

```
1313:'It's a Cheshire cat,' said the Duchess, 'and that's why. Pig!'
```

```
1319:'I didn't know that Cheshire cats always grinned; in fact, I didn't know
```

```
1433:was a little startled by seeing the Cheshire Cat sitting on a bough of a
```

```
1440:'Cheshire Puss,' she began, rather timidly, as she did not at all know
```

```
2068:'It's the Cheshire Cat: now I shall have somebody to talk to.'
```

```
2100:'It's a friend of mine--a Cheshire Cat,' said Alice: 'allow me to
```

```
2144:When she got back to the Cheshire Cat, she was surprised to find quite a
```

If you want more context to each match use the `-B` option to specify how many lines *before* a match should also be printed and `-A` for the number of lines *after* a match (output truncated after the first two matches to save space):

```
Gretchen:coolstuff thijscoenen$ grep -B 1 -A 2 -i "Cheshire" alice.txt
```

```
'It's a Cheshire cat,' said the Duchess, 'and that's why. Pig!'
```

```
She said the last word with such sudden violence that Alice quite
```

```
--
```

```
'I didn't know that Cheshire cats always grinned; in fact, I didn't know  
that cats COULD grin.'
```

```
--
```

Each match is separated by a line with two dashes --. If you are only interested in knowing whether a file contains a match you can use the `-H` option to print the filename and the matching text:

```
Gretchen:coolstuff thijscoenen$ grep -Hi "Cheshire" *.txt
```

```
alice.txt
```

```
Gretchen:coolstuff thijscoenen$
```

`grep` is a very powerful tool that can also match patterns, it is useful to read through the man pages to see what more it can do for you.

### 1.2.9 Glob Patterns

The BASH shell supports filename expansion, that is it will expand special wildcard symbols and match them to filenames. This is also called *globbing*, and the pattern that is matched *glob pattern*. I will first explain the parts that make up a glob pattern, then show you how you use them with BASH.

\* Matches any text in filename.

? Matches zero or one unspecified character.

[ABC] Allow one character from a specified set of characters to be matched, in this case the characters A, B and C make up that set. Will match a single character only, not ABC.

[0-9] Will match a number in the range 0 through 9.

[a-z] Match any lower case character in the range a through z.

[A-Z] Match any upper case character in the range A through Z.

[A-Za-z] Match any upper case character in the range A through Z or any lower case character in the range a through z.

\ Escape character, use this if you should match one of the characters that have a special meaning in a glob pattern. E.g. if you want to match the \*, your glob pattern should include \\* so that the \* itself matched and not expanded.

! This will negate a pattern. E.g. if you want to match anything not a lower case character use ![a-z].

Glob pattern can then be used in combination with other programs, what follows is a simple example that shows how `ls` can be combined with some glob patterns to look for some specific files.



```
Gretchen:glob-demonstration thijscoenen$ ls
101.txt          604.txt          ABC.txt
123.txt          709.dat          ABC123.txt
Gretchen:glob-demonstration thijscoenen$ ls [0-9]*
101.txt          123.txt          604.txt          709.dat
Gretchen:glob-demonstration thijscoenen$ ls [!0-9]*
ABC.txt          ABC123.txt
Gretchen:glob-demonstration thijscoenen$ ls [5-9]*
604.txt          709.dat
Gretchen:glob-demonstration thijscoenen$ ls *dat
709.dat
Gretchen:glob-demonstration thijscoenen$ ls 1[2468]*
123.txt
Gretchen:glob-demonstration thijscoenen$
```

### 1.2.10 Finding Files

A common problem is, how do I find some specific file? The naive approach, use `ls` and move around the file system until you find what you were looking for, is time consuming and . Beyond the glob patterns that your shell provides (see Section 1.2.9) the UNIX `find` command is useful. With `find` you can look through a directory hierarchy for filenames matching a certain pattern or certain properties. You can then run some commands for each file that was found. At its simplest you can use `find` to look for a certain file called `needle.txt` in this example:

```
Gretchen:find thijscoenen$ find . -name "needle.txt"
./haystack/hay/hay/hay/needle.txt
./haystack/straw/needle.txt
Gretchen:find thijscoenen$
```

As you can see two files were found, but what if we know that the file we need contains the word `gold`? It turns out you can run other programs on each found file using the `-exec` option of `find`. In the following example `grep` is used to check whether a match contains `gold`:

```
Gretchen:find thijscoenen$ find . -name "needle.txt" -exec grep -Hi "gold" {} \;
./haystack/straw/needle.txt
Gretchen:find thijscoenen$
```

It is now clear that the file we needed was `./haystack/straw/needle.txt`. In the commands that follow the `-exec` each occurrence of `{}` is replaced with the matching file name and each command must be followed by `\;`. The `grep` command was explained in Section 1.2.8.

### 1.2.11 gzip and tar

To conserve disk space it is sometimes a good idea to pack files more densely, there are several tools to do so. If you have used Windows so called zip files may be familiar. Although utilities to handle zip files exist (`zip` and `unzip`) on UNIX, you are more likely to come across `.gz` or `.bz2` files. Below I give examples of how to deflate and reinflate a file using `gzip`, note that I use a so-called pipe `|` and `grep` to only show filenames containing `alice`. Pipes are a way of sending output from one program to another, which are explained in Section 1.2.14. The following example shows that the Alice in Wonderland story can be packed to only about 36 % of its original size:

```
Gretchen:coolstuff thijscoenen$ ls -l | grep alice
-rw-r--r-- 1 thijscoenen staff 147731 22 jul 23:27 alice.txt
Gretchen:coolstuff thijscoenen$ gzip alice.txt
Gretchen:coolstuff thijscoenen$ ls -l | grep alice
```

```
-rw-r--r-- 1 thijscoenen staff 53596 22 jul 23:27 alice.txt.gz
Gretchen:coolstuff thijscoenen$ gunzip alice.txt.gz
Gretchen:coolstuff thijscoenen$ ls -l | grep alice
-rw-r--r-- 1 thijscoenen staff 147731 22 jul 23:27 alice.txt
Gretchen:coolstuff thijscoenen$
```

The `bzip2` and `bunzip2` programs work similarly and may achieve a slightly higher compression than `gzip`, but the latter is used more often.

When you download software or data you will come across so-called tarballs, which are files that can encapsulate many separate files or even directory full hierarchies. The name tarball derives from the `tar` utility that can create or unpack them. The name `tar` itself derives from tape archive (archiving to tape was and still is cheaper than archiving to disk). The example that follows shows how you create (`-c` option) a tarball that is compressed using `gzip` (`-z` option) how you can list its contents (`-t` option) and how you can extract it again (`-x` option):

```
Gretchen:coolstuff thijscoenen$ ls
Whatever      alice.txt      hello          noperm
Gretchen:coolstuff thijscoenen$ tar -czf cool.tar.gz *
Gretchen:coolstuff thijscoenen$ tar -tzf cool.tar.gz
Whatever
alice.txt
hello
noperm
Gretchen:coolstuff thijscoenen$ rm Whatever alice.txt noperm hello
Gretchen:coolstuff thijscoenen$ ls
cool.tar.gz
Gretchen:coolstuff thijscoenen$ tar -xzf cool.tar.gz
Gretchen:coolstuff thijscoenen$ ls
Whatever      cool.tar.gz    noperm
alice.txt      hello
Gretchen:coolstuff thijscoenen$ rm cool.tar.gz
Gretchen:coolstuff thijscoenen$
```

Note that some tarballs are compressed using `bzip2` in which case you can use the `-j` option instead of `-z`.

### 1.2.12 Permissions Continued

In Section 1.1.1 users and their permissions were introduced. UNIX has several commands that let you manage file permissions and ownership. UNIX also allows a user to temporarily acquire more privileges. We will start by demonstrating how file permissions can be changed. The owner of a file and a user with root privileges can change its permissions. The first step is to check the permissions, run for example `ls -l` to get a file listing that shows the permissions in the first column of the output.

```
Gretchen:chmod-demo thijscoenen$ ls -l
total 16
-rw-r--r-- 1 thijscoenen staff 68 27 aug 22:02 demo.py
----r--r-- 1 thijscoenen staff 15 27 aug 22:00 nopriv.txt
drwxr-xr-x 2 thijscoenen staff 68 27 aug 22:12 subdir
```

The first column can be further split, taking `demo.py` as an example we get the file type `-`, owner permissions `rw-`, group permissions `r--` and finally the permissions for everyone else `r--`. The file type can be: `-` for a file, `d` for a directory and `l` for a “symbolic link”. In our example we are dealing with a normal file. The user permissions are listed in the order read, write and execute using

the abbreviations `rw`. If a permission is not granted the permission will be replaced with a dash. In our example the owner of the file has read and write permissions but no execute permissions: `rw-`. The group permissions in the example are only read permissions: `r--`. All other users are also only able to read the file: `r--`. The second column in the output of `ls -l` is not interesting for our current purposes, but the third is: it shows the owner of the file. In this case that user is `thijscoenen`. The following column shows the group that the file is assigned to, in this case the group `staff`. The next column shows the file size in bytes, followed by a column that shows the last time the file was changed. The last column shows the actual file or directory name.

To change the permissions of the user, group and all other users the file owner (or a sufficiently privileged user like `root`) can use `chmod`. This command takes as options a string like `ugo+rw` that specifies which user should gain or lose which permissions and finally as arguments the files for which the permissions should be changed. The options should be read as: the owner, the group and all others (`ugo`), gain (+) the permissions to read, write and execute (`rw`). If we want to make the `demo.py` file executable for its owner, the following command will do the trick:

```
Gretchen:chmod-demo thijscoenen$ chmod u+x demo.py
Gretchen:chmod-demo thijscoenen$ ls -l
total 16
-rwxr--r-- 1 thijscoenen staff 82 27 aug 22:29 demo.py
----r--r-- 1 thijscoenen staff 15 27 aug 22:00 nopriv.txt
drwxr-xr-x 2 thijscoenen staff 68 27 aug 22:12 subdir
Gretchen:chmod-demo thijscoenen$
```

If we want to disallow other users to see what is in the `subdir` directory we can remove (-) the execute permissions on that directory for its group and all other users:

```
Gretchen:chmod-demo thijscoenen$ chmod go-x subdir/
Gretchen:chmod-demo thijscoenen$ ls -l
total 16
-rwxr--r-- 1 thijscoenen staff 82 27 aug 22:29 demo.py
----r--r-- 1 thijscoenen staff 15 27 aug 22:00 nopriv.txt
drwxr--r-- 2 thijscoenen staff 68 27 aug 22:12 subdir
Gretchen:chmod-demo thijscoenen$
```

You can also specify several files by just typing their names or by using the appropriate glob patterns. To recursively change the permissions of some directory hierarchy use the `-R` option, for example to change the permissions on `subdir` and everything it contains to be completely open to everyone: `chmod -R ugo=rwx subdir`. In this last command we use `=` to just set the permissions. Old examples explaining `chmod` may show numerical permissions, but they are not necessary to manage the permissions and will not be treated here.

To change the owner of a file use the `chown` command. Using requires that you are either the owner or are otherwise sufficiently privileged. To change the owner of `demo.py` from the examples to `root` use the following command:

```
Gretchen:chmod-demo thijscoenen$ chown root:staff demo.py
chown: demo.py: Operation not permitted
Gretchen:chmod-demo thijscoenen$ sudo chown root:staff demo.py
Password:
Gretchen:chmod-demo thijscoenen$ ls -l
total 16
-rwxr--r-- 1 root          staff 82 27 aug 22:29 demo.py
----r--r-- 1 thijscoenen staff 15 27 aug 22:00 nopriv.txt
drwxrwxrwx 2 thijscoenen staff 68 27 aug 22:12 subdir
```

As you can see there is a small hitch, the user `thijscoenen` cannot change the ownership to `root` and group `staff`, because that user is not privileged enough. The second time `chmod` is preceded with

`sudo`, which means execute the command that follows as a super user (root on most systems). To prove that you in fact are allowed to acquire root privileges you are asked to enter the password that belongs to that user. After acquiring the privileges `chown` to root becomes possible. While this example is a bit silly `sudo` is very important when installing software. A lot of software is installed system wide and that means somewhere in the system directories of UNIX. Those directories are not writable for normal user, a problem that is solved with `sudo`. Note: if you are using a university computer you will most likely not have the ability to use `sudo`, because most system administrators will not give out root privileges to just anyone — on your own Mac or Linux machine that should be no problem.

### 1.2.13 Environment Variables

Some of the settings of your BASH shell are put in so-called *environment variables*. A *variable* is a name (or an identifier) associated with some stored value that can be changed. We already came across an environment variable called `$PATH`, its value is a list of directories that are searched for executables. The BASH `echo $PATH` command will echo (show) the value of the variable `$PATH`, the `$` is needed to signify that the following name is a variable (the example below shows what goes wrong without the `$`):

```
Gretchen:~ thijscoenen$ echo $PATH
/Library/Frameworks/Python.framework/Versions/2.7/bin:/usr/bin:/bin:/usr/sbin:/sbin:/usr/local/bin:/usr/texbin:/usr/X11/bin
Gretchen:~ thijscoenen$ echo PATH
PATH
Gretchen:~ thijscoenen$
```

As you can see `$PATH` contains multiple directories separated with a colon `:`. When you change `$PATH` you should know that the directories are searched for a executable program in the same order as they are specified in `$PATH`.

To set a variable to a certain value use the `export` command, e.g. to set `$GREETING` to Good morning use the following command:

```
Gretchen:~ thijscoenen$ export GREETING="Good morning"
Gretchen:~ thijscoenen$ echo $GREETING
Good morning
Gretchen:~ thijscoenen$
```

This variable will be defined as long as your BASH session exists. If you are running several BASH sessions, the variable will only be defined for that one session where it was defined. To define a variable for all future BASH session you should edit the settings file for your BASH sessions. This file is probably called `.bashrc` (Linux) or `.bash_profile` (Mac OS X) in your home directory. Be careful editing these files, getting it wrong may make your shell inoperable.

If you want to know which environment variables are defined in your shell and what their values are you can use `env`. Running this command will result in a list of environment variables and their values. Because this is a long list I suggest you run this command on your own computer.

### 1.2.14 UNIX Pipes

So far we have treated a number of simple commands and utilities available on the command line. The power of the UNIX command line partially derives from the ability to connect different commands. The output of one program can be sent to another program using a so-called pipe. A pipe is signified with the `|` character. E.g. to use `grep` to look for a file called `.bash_profile` in my home directory:

```
Gretchen:~ thijscoenen$ ls -la | grep ".bash_profile"
-rw-r--r--  1 thijscoenen  staff      573 22 jul 23:38 .bash_profile
Gretchen:~ thijscoenen$
```

This may be a very simple example but the principle generalizes easily. The next example shows how you can search through a directory with some type of pulsar data files, looking for the dispersion measure, sorting by dispersion measure and then looking at only the first 10 lines:

```
Gretchen:00000143 thijscoenen$ ls *.inf | wc -l      313
Gretchen:00000143 thijscoenen$ cat *.inf | grep "Dispe" | sort -n -k 6 | head -n 10
Dispersion measure (cm-3 pc)      = 2.6
Dispersion measure (cm-3 pc)      = 2.65
Dispersion measure (cm-3 pc)      = 2.7
Dispersion measure (cm-3 pc)      = 2.75
Dispersion measure (cm-3 pc)      = 2.8
Dispersion measure (cm-3 pc)      = 2.85
Dispersion measure (cm-3 pc)      = 2.9
Dispersion measure (cm-3 pc)      = 2.95
Dispersion measure (cm-3 pc)      = 3
Dispersion measure (cm-3 pc)      = 3.05
Gretchen:00000143 thijscoenen$
```

Understanding how you can combine different UNIX command line utilities can help you automate boring tasks. This document is not about BASH programming so I will stop here now.

### 1.2.15 Output Redirects

The output of programs that shows up in your shell was written to “standard out” or “standard error” (a program should use the former for normal output while the latter can be used for errors or warnings). If you want these messages in a text file, you could copy the text from your shell into a text editor and save a file, but that is a lot of work. Instead you can just redirect standard out to a text file with a `>` sign:

```
Gretchen:coolstuff thijscoenen$ ls -l > blah.txt
Gretchen:coolstuff thijscoenen$ cat blah.txt
total 312
-rw-r--r--  1 thijscoenen  staff      0 15 aug 16:35 Whatever
-rw-r--r--  1 thijscoenen  staff 147731 22 jul 23:27 alice.txt
-rw-r--r--  1 thijscoenen  staff      0 28 aug 01:31 blah.txt
-rwxr--r--  1 thijscoenen  staff    43  7 aug 18:02 hello
-rw-r--r--  1 thijscoenen  staff    50 11 aug 14:35 noperm
Gretchen:coolstuff thijscoenen$
```

If you want to capture both standard out and standard error in a file you can add `2>&1` to the end of the redirect command:

```
Gretchen:~ thijscoenen$ python divide_by_zero.py
Hi I'll divide by zero!
The result is
Traceback (most recent call last):
  File "divide_by_zero.py", line 2, in <module>
    print "The result is", 1/0
ZeroDivisionError: integer division or modulo by zero
Gretchen:~ thijscoenen$ python divide_by_zero.py > output1.txt
Traceback (most recent call last):
```

```
File "divide_by_zero.py", line 2, in <module>
    print "The result is", 1/0
ZeroDivisionError: integer division or modulo by zero
Gretchen:~ thijscoenen$ cat output1.txt
Hi I'll divide by zero!
The result is
Gretchen:~ thijscoenen$ python divide_by_zero.py > output2.txt 2>&1
Gretchen:~ thijscoenen$ cat output2.txt
Hi I'll divide by zero!
The result is
Traceback (most recent call last):
  File "divide_by_zero.py", line 2, in <module>
    print "The result is", 1/0
ZeroDivisionError: integer division or modulo by zero
Gretchen:~ thijscoenen$
```

In this example you can see that Python sends the text associated with an `ZeroDivisionError` to standard error and the rest of the messages to standard out. The file `output1.txt` only contains the text that went to standard out, while the file `output2.txt` contains both the text for standard out and standard error. In case you do not want to overwrite an existing file, but append to the new output to it replace the `>` with `>>`.

# Part II

## Assignments

## Chapter 2

# Finding Pulsars from the command line

### 2.1 Introduction

Pulsar surveys in radio produce large numbers of potential pulsars, these then have to be inspected. For this exercise you are given a large number of potential pulsars. For each potential pulsar you will find a `.bestprof`-file and a `.png`-file. The former is a text file with meta data while the latter is a diagnostic plot. Below I will use detection to describe one of these potential pulsar signals (whether a real pulsar or not). The data set should contain a pulsar that you have to find using command line tools.

The data you need is available from the <http://www.astroprog.nl> website.

#### 2.1.1 Grading

While this exercise does not count towards your final grade, it is required to be handed in.

#### 2.1.2 The assignments

- Start a report to be handed in as a A4-sized PDF file. Make sure you include your name and student number on the first page of your report.
- Download the files `*.tar.gz` (containing a number of potential pulsar candidates) and copy it to a directory where you can work.
- Unpack the tarballs, you will now be presented with a number of sub directories containing the images and metadata files.
- Open one of these `.png` images from the command line. If you don't know which command to use to open these images, perform a web search to find out which command will do that. (Add the command you used to the report.)
- Now that you have seen the plot, open the metadata (hint: the `bestprof` files are text files). Looking through this file you will see that the reduced chi-squared is an entry in the header, we will need these later. How many entries are there in the header? (Add the number to your report.)
- How many candidates are there in the complete data set? (Add the answer to your report.)
- Add the picture of the brightest profile to your report (the brightest candidate will have the highest reduced chi squared). You should use the meta data to find the brightest candidate. Also add to the report how you found the brightest one.



- Now add the 83th brightest profile you report and explain how you found it. (The reader chapter about BASH contains a hint as to how you could do this.)